

Optimization and Numerical Methods Solutions

Thom Volker

28-10-2022

Table of contents

Preface	3
1 Introduction	4
2 Motivating Problems	5
2.1 Exercises (2.7 in the notes)	5
3 Basic tools	7
3.1 Exercises (3.7 in the book)	7
4 From non-iterative to iterative procedures	20
5 Least squares	21
6 Iteration-based Function Optimization	22
6.1 Exercises (6.5 in the notes)	22
References	41

Preface

This project has two purposes. First, it is an attempt to organize my solutions to the course Optimization and Numerical Methods in a structured way. Second, it provides a justification to try and learn Quarto.

1 Introduction

No exercises.

2 Motivating Problems

Chapter 2 on motivating problems is the first chapter that actually entails exercises.

2.1 Exercises (2.7 in the notes)

1. Consider the multinomial likelihood in Equation 2.1 for a model (for a two-way contingency table) assuming independence. Can you simplify the likelihood?

$$\sum_{j=1}^R \sum_{k=1}^C n_{jk} \ln(\pi_{jk}) \quad \sum_{j=1}^R \sum_{k=1}^C \pi_{jk} = 1 \quad (2.1)$$

Solution

$$\begin{aligned} \ell(\pi) &= \sum_{j=1}^R \sum_{k=1}^C n_{jk} \ln(\pi_{jk}) \\ &= \sum_{j=1}^R \sum_{k=1}^C n_{jk} \ln(\pi_{j+} \cdot \pi_{+k}) \\ &= \sum_{j=1}^R \sum_{k=1}^C n_{jk} \ln \pi_{j+} + n_{jk} \ln \pi_{+k} \\ &= \sum_{j=1}^R n_{j+} \ln \pi_{j+} + \sum_{k=1}^C n_{+k} \ln \pi_{+k} \end{aligned} \quad (2.2)$$

2. In a mixed model, optimization is carried out using the marginal likelihood (the likelihood with the random effects integrated out). Define the marginal likelihood for the one-way random effects ANOVA model.

One-way random effects ANOVA with group-specific effects $\mu_j \sim \mathcal{N}(0, \sigma_\mu^2)$, and

$$y_{ij} = \beta + \mu_j + \epsilon_{ij},$$

with $\epsilon \sim \mathcal{N}(0, \sigma_\epsilon^2)$, with a groups indexed j , and n_j individuals in every group.

Solution

So, the likelihood consists of two components. For the individuals within each group, we have

$$\prod_{i=1}^{n_j} \frac{1}{\sqrt{2\pi\sigma_\epsilon^2}} \exp\left(-\frac{(y_{ij} - \beta - \mu_j)^2}{2\sigma_\epsilon^2}\right),$$

whereas for the groups themselves, we have

$$\prod_{j=1}^a \frac{1}{\sqrt{2\pi\sigma_\mu^2}} \exp\left(-\frac{\mu_j^2}{2\sigma_\mu^2}\right).$$

Combining these components, and integrating out the random effects, we obtain the marginal likelihood

$$\prod_{j=1}^a \int \prod_{i=1}^{n_j} \frac{1}{\sqrt{2\pi\sigma_\epsilon^2}} \exp\left(-\frac{(y_{ij} - \beta - \mu_j)^2}{2\sigma_\epsilon^2}\right) \frac{1}{\sqrt{2\pi\sigma_\mu^2}} \exp\left(-\frac{\mu_j^2}{2\sigma_\mu^2}\right) d\mu_j.$$

3. Suppose you do a simple linear regression analysis using a t_ν -distribution for the residuals (density: $f_\nu(y) = C\sqrt{\lambda}\left(1 + \frac{\lambda(y-\mu)^2}{\nu}\right)^{-\frac{\nu+1}{2}}$ where μ is the mean (for $\nu > 1$), λ is a scale parameter and C is a normalizing constraint that does not depend on μ or λ). Define the (log-)likelihood for n observations (y_i, x_i) , such that $\mu_i = \beta_0 + \beta_1 x_i$.

Solution

$$L(\beta) = \prod_{i=1}^n C\sqrt{\lambda} \left(1 + \frac{\lambda(y_i - \beta_0 - \beta_1 x_i)^2}{\nu}\right)^{-\frac{\nu+1}{2}},$$

$$\ell(\beta) = N \ln C + \frac{N}{2} \ln \lambda - \sum_{i=1}^n \frac{\nu+1}{2} \ln \left(1 + \frac{\lambda(y_i - \beta_0 - \beta_1 x_i)^2}{\nu}\right)$$

3 Basic tools

Chapter 3 introduces basic tools for optimization problems, such as Taylor Series Expansion, and introduces the exponential family.

3.1 Exercises (3.7 in the book)

1. Consider $f(x) = \frac{e^x}{1+e^x}$. Derive the third-order Taylor series expansion of this function at $x = 0$, and make a graph with the function and the third-order Taylor series expansion at $x = 0$.

Solution

$$\begin{aligned} f(x) &= \frac{e^x}{1+e^x} \\ f'(x) &= \frac{e^x(1+e^x)}{(1+e^x)^2} - \frac{e^{2x}}{(1+e^x)^2} = \frac{e^x}{(1+e^x)^2} \\ f''(x) &= \frac{e^x(1+e^x)^2 - e^{2x}2(1+e^x)e^x}{(1+e^x)^4} \\ &= \frac{e^x(1+e^x)^2 - 2e^{2x}}{(1+e^x)^3} \\ &= \frac{e^x - e^{2x}}{(1+e^x)^3} \\ f'''(x) &= \frac{(e^x - 2e^{2x})(1+e^x)^3 - (e^x - e^{2x})3(1+e^x)^2e^x}{(1+e^x)^6} \\ &= \frac{e^x - 2e^{2x} + e^{2x} - 2e^{3x} - 3e^{2x} + 3e^{3x}}{(1+e^x)^4} \\ &= \frac{e^x - 4e^{2x} + e^{3x}}{(1+e^x)^4}, \end{aligned}$$

using Taylor's theorem, we get

$$f(x) \approx \sum_{k=0}^n \frac{f^{(k)}(x_0)}{k!} (x - x_0)^k$$


$$= \frac{1}{2} + \frac{1}{4}x + 0 - \frac{1}{48}x^3.$$

```
library(ggplot2)
fx <- function(x) exp(x) / (1 + exp(x))
fx1 <- function(x) exp(x) / (1 + exp(x))^2
fx2 <- function(x) (exp(x) - exp(2*x)) / (1 + exp(x))^3
fx3 <- function(x) (exp(x) - 4*exp(2*x) + exp(3*x)) / (1 + exp(x))^4

taylor <- function(x, root) {
  fx(root) + fx1(root) * (x - root) + fx2(root) / 2 * (x - root)^2 + fx3(root) / 6 * (x -
}

ggplot() +
  geom_function(fun = fx) +
  geom_function(fun = taylor, args = list(root = 0), col = "orange") +
  xlim(-3, 3) +
  theme_minimal() +
  labs(x = "X", y = expression(italic(f(X))),
       title = "Third-order Taylor Series Expansion")
```


Third-order Taylor Series Expansion



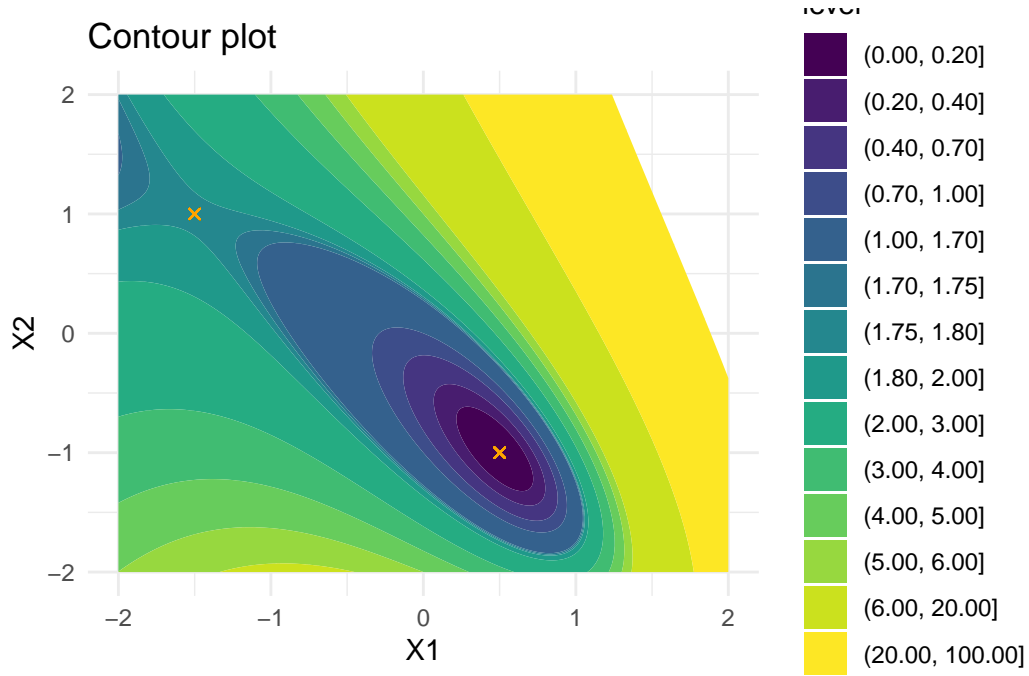
2. Consider the function: $f(x) = e^{x_1}(4x_1^2 + 2x_2^2 + 4x_1x_2 + 2x_2 + 1)$. Make a contour plot of this function (let both axes run from -2 to 2) at function values 0.2, 0.4, 0.7, 1, 1.7, 1.75, 1.8, 2, 3, 4, 5, 6, 20. Derive the second-order Taylor series at $x = (0.5, -1)'$ and $x = (-0.75, 1)'$.

Solution

Contour plot

```
fx12 <- function(x1, x2) {
  exp(x1) * (4*x1^2 + 2*x2^2 + 4*x1*x2 + 2*x2 + 1)
}

expand.grid(x1 = -200:200/100,
            x2 = -200:200/100) |>
  dplyr::mutate(z = fx12(x1, x2)) |>
  ggplot(aes(x = x1, y = x2, z = z)) +
  stat_contour_filled(breaks = c(0, 0.2, 0.4, 0.7, 1, 1.7, 1.75, 1.8, 2, 3, 4, 5, 6, 20, 100)) +
  geom_point(aes(x = 0.5, y = -1), col = "orange", shape = "cross") +
  geom_point(aes(x = -1.5, y = 1), col = "orange", shape = "cross") +
  theme_minimal() +
  labs(x = "X1", y = "X2",
       title = "Contour plot")
```



The second-order Taylor expansion uses the first and second partial derivatives of the function $f(x)$.

$$\begin{aligned}
 f(x) &= e^{x_1}(4e_1^2 + 2x_2^2 + 4x_1x_2 + 2x_2 + 1), \\
 \frac{\partial f}{\partial x_1} &= f(x) + e^{x_1}(8x_1 + 4x_2), \\
 \frac{\partial f}{\partial x_2} &= e^{x_1}(4x_2 + 4x_1 + 2), \\
 \frac{\partial^2 f}{\partial x_1^2} &= f(x) + 2e^{x_1}(8x_1 + 4x_2) + 8e^{x_1}, \\
 \frac{\partial^2 f}{\partial x_2^2} &= 4e^{x_1}, \\
 \frac{\partial^2 f}{\partial x_1 \partial x_2} &= 4e^{x_1} + e^{x_1}(4x_2 + 4x_1 + 2).
 \end{aligned}$$

Accordingly, the Gradient $\nabla f(x)$ is defined as

$$\nabla f(x) = \begin{pmatrix} f(x) + e^{x_1}(8x_1 + 4x_2) \\ e^{x_1}(4x_2 + 4x_1 + 2) \end{pmatrix},$$

and the Hessian $\nabla^2 f(x)$ is defined as

$$\nabla^2 f(x) = \begin{pmatrix} f(x) + 2e^{x_1}(8x_1 + 4x_2) + 8e^{x_1} & 4e^{x_1} + e^{x_1}(4x_2 + 4x_1 + 2) \\ 4e^{x_1} + e^{x_1}(4x_2 + 4x_1 + 2) & 4e^{x_1} \end{pmatrix}.$$

Moreover, the second-order Taylor series at $x = (0.5, -1)'$ and $x = (-0.75, 1)'$ is defined as

$$\begin{aligned} \nabla f((0.5, -1)) &= \begin{pmatrix} 0 \\ 0 \end{pmatrix} \\ \nabla^2 f((0.5, -1)) &= \begin{pmatrix} 13.19 & 6.59 \\ 6.59 & 6.59 \end{pmatrix}, \end{aligned}$$

and

$$\begin{aligned} \nabla f((-1.5, 1)) &= \begin{pmatrix} 0 \\ 0 \end{pmatrix} \\ \nabla^2 f((-1.5, 1)) &= \begin{pmatrix} 0 & 0.89 \\ 0.89 & 0.89 \end{pmatrix}. \end{aligned}$$

As can be seen in the contour plot, the first point is a minimum, while the second point is a saddle point.

3. Consider the likelihood function

$$L = \prod_{i=1}^N \frac{e^{(\alpha + \beta x_i)y_i}}{1 + e^{(\alpha + \beta x_i)}}.$$

derive the log-likelihood function, the gradient vector for the parameter vector $\theta = (\alpha, \beta)$ and the Hessian matrix for the parameter vector θ .

Solution

The log-likelihood is defined as

$$\ell = \sum_{i=1}^N (\alpha + \beta x_i)y_i - \log(1 + e^{(\alpha + \beta x_i)}),$$

differentiation with respect to α yields

$$\frac{\partial \ell}{\partial \alpha} = \sum_{i=1}^N y_i - \frac{e^{(\alpha + \beta x_i)}}{1 + e^{(\alpha + \beta x_i)}} = \sum_{i=1}^N y_i - \pi_i,$$

differentiation with respect to β yields

$$\frac{\partial \ell}{\partial \beta} = \sum_{i=1}^N y_i x_i - x_i \frac{e^{(\alpha + \beta x_i)}}{1 + e^{(\alpha + \beta x_i)}} = \sum_{i=1}^N x_i (y_i - \pi_i).$$

Accordingly, the gradient is defined as

$$\nabla \ell = \begin{pmatrix} \sum_{i=1}^N y_i - \pi_i \\ \sum_{i=1}^N x_i (y_i - \pi_i) \end{pmatrix}.$$

The second partial derivatives are defined as

$$\begin{aligned} \frac{\partial^2 \ell}{\partial \alpha^2} &= \sum_{i=1}^N -\frac{e^{(\alpha + \beta x_i)}(1 + e^{(\alpha + \beta x_i)}) - e^{(\alpha + \beta x_i)}e^{(\alpha + \beta x_i)}}{(1 + e^{(\alpha + \beta x_i)})^2} \\ &= -\sum_{i=1}^N \frac{e^{(\alpha + \beta x_i)}}{1 + e^{(\alpha + \beta x_i)}} - \frac{(e^{(\alpha + \beta x_i)})^2}{(1 + e^{(\alpha + \beta x_i)})^2} \\ &= -\sum_{i=1}^N \pi_i(1 - \pi_i), \\ \frac{\partial^2 \ell}{\partial \alpha^2} &= \sum_{i=1}^N -x_i^2 \frac{e^{(\alpha + \beta x_i)}(1 + e^{(\alpha + \beta x_i)}) - e^{(\alpha + \beta x_i)}e^{(\alpha + \beta x_i)}}{(1 + e^{(\alpha + \beta x_i)})^2} \\ &= -\sum_{i=1}^N x_i^2 \pi_i(1 - \pi_i), \\ \frac{\partial^2 \ell}{\partial \alpha \partial \beta} &= \sum_{i=1}^N -x_i \frac{e^{(\alpha + \beta x_i)}(1 + e^{(\alpha + \beta x_i)}) - e^{(\alpha + \beta x_i)}e^{(\alpha + \beta x_i)}}{(1 + e^{(\alpha + \beta x_i)})^2} \\ &= -\sum_{i=1}^N x_i \pi_i(1 - \pi_i). \end{aligned}$$

Hence, the Hessian $\nabla^2 \ell$ is defined as

$$\nabla^2 \ell = \begin{pmatrix} -\sum_{i=1}^N \pi_i(1 - \pi_i) & -\sum_{i=1}^N x_i \pi_i(1 - \pi_i) \\ -\sum_{i=1}^N x_i \pi_i(1 - \pi_i) & -\sum_{i=1}^N x_i^2 \pi_i(1 - \pi_i) \end{pmatrix}.$$

4. Take the Weibull density

$$p(y) = \varphi \rho y^{\rho-1} e^{-\varphi y^\rho}.$$

Derive the second-order Taylor series expansion of $p(y)$ about $y = 1$.

Solution

$$\begin{aligned}\frac{\partial}{\partial y} [\varphi \rho y^{\rho-1} e^{-\varphi y^\rho}] &= \varphi \rho \left((\rho-1) y^{\rho-2} e^{-\varphi y^\rho} - \varphi \rho y^{2\rho-2} e^{-\varphi y^\rho} \right) \\ &= \varphi \rho e^{-\varphi y^\rho} y^{\rho-2} (\rho-1 - \varphi \rho y^\rho), \\ \frac{\partial^2}{\partial y^2} [\varphi \rho y^{\rho-1} e^{-\varphi y^\rho}] &= \varphi \rho \left[\frac{\partial}{\partial y} \rho \left(e^{-\varphi y^\rho} y^{\rho-2} \right) - \frac{\partial}{\partial y} \left(e^{-\varphi y^\rho} y^{\rho-2} \right) - \frac{\partial}{\partial y} \varphi \rho \left(e^{-\varphi y^\rho} y^{2\rho-2} \right) \right] \\ &= \varphi \rho e^{-\varphi y^\rho} y^{\rho-3} \left((\rho-1)(\rho-2 - \varphi \rho y^\rho) - \varphi \rho y^{2\rho-2} (2\rho-2 - \varphi \rho y^\rho) \right)\end{aligned}$$

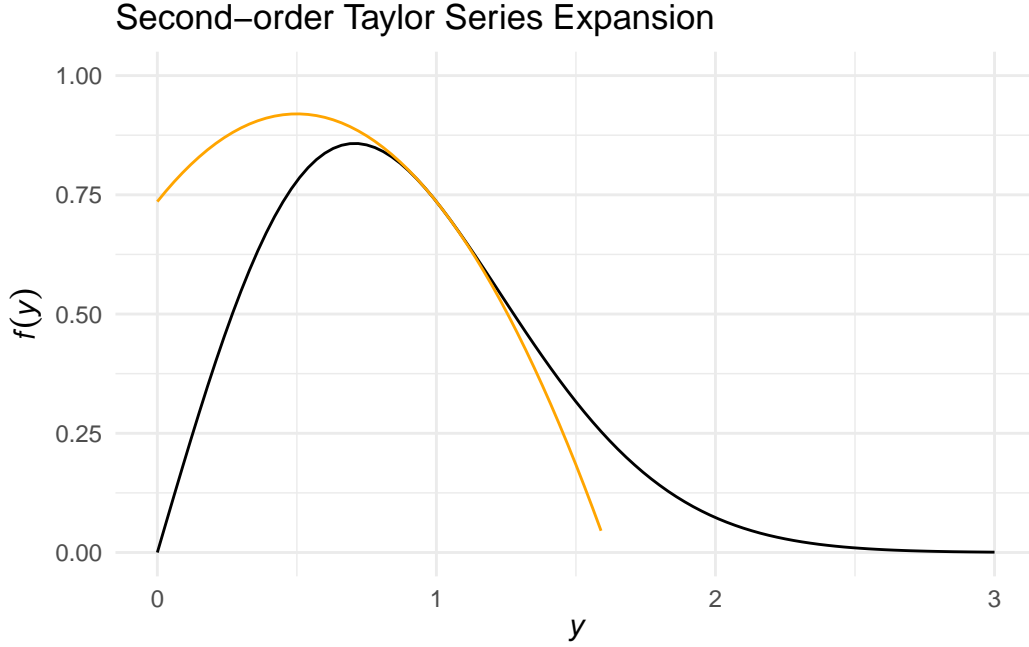
```
fx <- function(phi, rho, y) {
  e <- exp(-phi*y^rho)
  phi * rho * y^{rho-1} * e
}

fx1 <- function(phi, rho, y) {
  e <- exp(-phi*y^rho)
  phi*rho*e*y^{rho-2}*((rho-1) - phi*rho*y^rho)
}

fx2 <- function(phi, rho, y) {
  e <- exp(-phi*y^rho)
  phi*rho*e*y^{rho-3} * ((rho-1)*(rho-2-phi*rho*y^rho) - phi*rho*y^rho*(2*rho-2-phi*rho*y^rho))
}

taylor <- function(phi, rho, y, root) {
  fx(phi, rho, root) + fx1(phi, rho, root) * (y - root) + fx2(phi, rho, root)/2 * (y - root)^2
}

ggplot() +
  geom_function(fun = fx, args = list(phi = 1, rho = 2)) +
  geom_function(fun = taylor,
    args = list(phi = 1, rho = 2, root = 1),
    col = "orange") +
  lims(x = c(0, 3), y = c(0, 1)) +
  theme_minimal() +
  labs(x = expression(italic(y)), y = expression(italic(f(y))),
    title = "Second-order Taylor Series Expansion")
```



5. Consider the Weibull-based likelihood function:

$$L = \prod_{i=1}^n \rho y_i^{\rho-1} e^{(\alpha+\beta x_i)} e^{-(y_i^\rho e^{(\alpha+\beta x_i)})},$$

with y_i the outcome (time-to-event), x_i is a continuous covariate, and α and β are regression parameters. Derive the log-likelihood function for an i.i.d. sample of n observations (y_1, y_2, \dots, y_n) , the gradient of the log-likelihood function for the parameters (ρ, α, β) and the Hessian of the log-likelihood function for the parameter vector (ρ, α, β) .

Solution

The log-likelihood is defined as

$$\ell = \sum_{i=1}^n \log(\rho) + (\rho - 1) \log(y_i) + \alpha + \beta x_i - y_i^\rho e^{(\alpha+\beta x_i)}.$$

The first-order partial derivatives with respect to ρ, α, β are given by

$$\begin{aligned}\frac{\partial \ell}{\partial \rho} &= \sum_{i=1}^n \rho^{-1} + \log(y_i) - y_i^\rho e^{(\alpha + \beta x_i)} \log(y_i), \\ \frac{\partial \ell}{\partial \alpha} &= \sum_{i=1}^n 1 - y_i^\rho e^{(\alpha + \beta x_i)}, \\ \frac{\partial \ell}{\partial \beta} &= \sum_{i=1}^n x_i (1 - y_i^\rho e^{(\alpha + \beta x_i)}),\end{aligned}$$

such that the gradient is defined as

$$\nabla \ell = \begin{pmatrix} \sum_{i=1}^n \rho^{-1} + \log(y_i) - y_i^\rho e^{(\alpha + \beta x_i)} \log(y_i), \\ \sum_{i=1}^n 1 - y_i^\rho e^{(\alpha + \beta x_i)}, \\ \sum_{i=1}^n x_i (1 - y_i^\rho e^{(\alpha + \beta x_i)}), \end{pmatrix}.$$

Additionally, the second-order partial derivatives are defined by

$$\begin{aligned}\frac{\partial^2 \ell}{\partial \rho^2} &= \sum_{i=1}^n -\rho^{-2} - y_i^\rho e^{(\alpha + \beta x_i)} (\log(y_i))^2, \\ \frac{\partial^2 \ell}{\partial \alpha^2} &= \sum_{i=1}^n -y_i^\rho e^{(\alpha + \beta x_i)}, \\ \frac{\partial^2 \ell}{\partial \beta^2} &= \sum_{i=1}^n -x_i^2 y_i^\rho e^{(\alpha + \beta x_i)}, \\ \frac{\partial^2 \ell}{\partial \rho \partial \alpha} &= \sum_{i=1}^n -\log(y_i) y_i^\rho e^{(\alpha + \beta x_i)}, \\ \frac{\partial^2 \ell}{\partial \rho \partial \beta} &= \sum_{i=1}^n -x_i \log(y_i) y_i^\rho e^{(\alpha + \beta x_i)}, \\ \frac{\partial^2 \ell}{\partial \alpha \partial \beta} &= \sum_{i=1}^n -x_i y_i^\rho e^{(\alpha + \beta x_i)},\end{aligned}$$

such that the Hessian is defined as

$$\nabla^2 \ell(\rho, \alpha, \beta) = \begin{pmatrix} \sum_{i=1}^n -\rho^{-2} - y_i^\rho e^{(\alpha + \beta x_i)} (\log(y_i))^2 & \sum_{i=1}^n -\log(y_i) y_i^\rho e^{(\alpha + \beta x_i)} & \sum_{i=1}^n -x_i \log(y_i) y_i^\rho e^{(\alpha + \beta x_i)} \\ \sum_{i=1}^n -\log(y_i) y_i^\rho e^{(\alpha + \beta x_i)} & \sum_{i=1}^n -y_i^\rho e^{(\alpha + \beta x_i)} & \sum_{i=1}^n -x_i y_i^\rho e^{(\alpha + \beta x_i)} \\ \sum_{i=1}^n -x_i \log(y_i) y_i^\rho e^{(\alpha + \beta x_i)} & \sum_{i=1}^n -x_i y_i^\rho e^{(\alpha + \beta x_i)} & \sum_{i=1}^n -x_i^2 y_i^\rho e^{(\alpha + \beta x_i)} \end{pmatrix}.$$

6. Consider a logistic regression

$$\text{logit}[P(Y_i = 1|x_i)] = \alpha + \beta x_i,$$

and a small set of data

i	x_i	y_i
1	0.5	0
2	1.0	0
3	1.5	1
4	2.0	0
5	2.5	1

Construct the log-likelihood function and the gradient function.

Solution

Constructing the logit function requires an expression for $P(Y_i = 1|x_i)$, which is defined as follows.

$$\begin{aligned}\text{logit}[P(Y_i = 1|x_i)] &= \alpha + \beta x_i, \\ \log\left(\frac{P(Y_i = 1|x_i)}{1 - P(Y_i = 1|x_i)}\right) &= e^{(\alpha + \beta x_i)}, \\ P(Y_i = 1|x_i) &= e^{(\alpha + \beta x_i)} - e^{(\alpha + \beta x_i)}(P(Y_i = 1|x_i)), \\ 1 &= \frac{e^{(\alpha + \beta x_i)}}{P(Y_i = 1|x_i)} - e^{(\alpha + \beta x_i)}, \\ 1 + e^{(\alpha + \beta x_i)} &= \frac{e^{(\alpha + \beta x_i)}}{P(Y_i = 1|x_i)}, \\ P(Y_i = 1|x_i) &= \frac{e^{(\alpha + \beta x_i)}}{1 + e^{(\alpha + \beta x_i)}}.\end{aligned}$$

Plugging this into a binomial likelihood function yields

$$\begin{aligned}
L &= \prod_{i=1}^5 \pi_i^{y_i} (1 - \pi_i)^{(1-y_i)}, \\
\ell &= \sum_{i=1}^5 y_i \log \pi_i + (1 - y_i) \log(1 - \pi_i) \\
&= \sum_{i=1}^5 y_i \log \left(\frac{e^{(\alpha + \beta x_i)}}{1 + e^{(\alpha + \beta x_i)}} \right) + \log \left(\frac{1}{1 + e^{(\alpha + \beta x_i)}} \right) - y_i \log \left(\frac{1}{1 + e^{(\alpha + \beta x_i)}} \right) \\
&= \sum_{i=1}^5 y_i \log \left(\frac{e^{(\alpha + \beta x_i)}}{1 + e^{(\alpha + \beta x_i)}} \right) + \log \left(\frac{1}{1 + e^{(\alpha + \beta x_i)}} \right) \\
&= \sum_{i=1}^n y_i (\alpha + \beta x_i) - \log(1 + e^{(\alpha + \beta x_i)}).
\end{aligned}$$

Accordingly, we can define the Gradient as

$$\nabla \ell = \begin{pmatrix} \sum_{i=1}^5 y_i - \frac{e^{(\alpha + \beta x_i)}}{1 + e^{(\alpha + \beta x_i)}} = \sum_{i=1}^5 y_i - \pi_i \\ \sum_{i=1}^5 y_i x_i - x_i \frac{e^{(\alpha + \beta x_i)}}{1 + e^{(\alpha + \beta x_i)}} = \sum_{i=1}^5 x_i (y_i - \pi_i) \end{pmatrix}$$

Filling in the values for y yields

$$\begin{aligned}
\frac{\partial \ell}{\partial \alpha} &= 2 - \sum_{i=1}^5 \pi_i, \\
\frac{\partial \ell}{\partial \beta} &= 4 - \sum_{i=1}^5 x_i \pi_i.
\end{aligned}$$

```

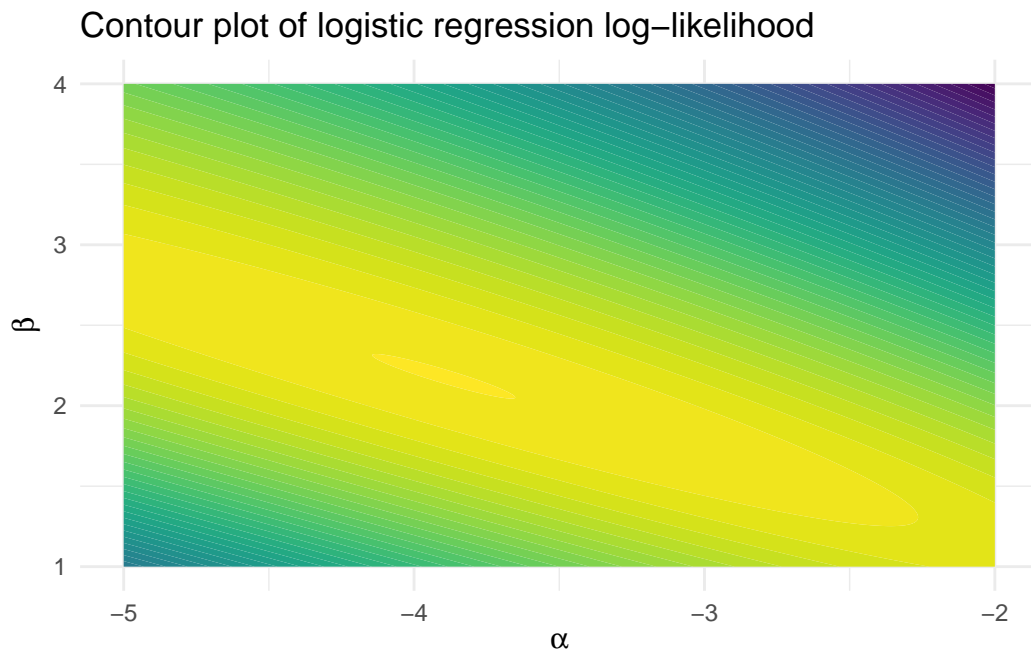
ell <- function(x, y, alpha, beta) {
  sum(y*(alpha + beta*x) - log(1 + exp(alpha + beta*x)))
}

x <- c(0.5, 1, 1.5, 2, 2.5)
y <- c(0, 0, 1, 0, 1)

expand.grid(alpha = -5000:-2000/1000,
            beta = 1000:4000/1000) |>
  dplyr::mutate(l = purrr::map2_dbl(alpha, beta, ~ell(x, y, .x, .y))) |>
  ggplot(aes(x = alpha, y = beta, z = l)) +
  stat_contour_filled(bins = 50, show.legend = FALSE) +
  theme_minimal() +

```

```
labs(x = expression(alpha),
     y = expression(beta),
     title = "Contour plot of logistic regression log-likelihood")
```



7. Consider $f(x_1, x_2, x_3) = (x_1 - 1)^4 + (x_2 - 3)^2 + 4(x_3 + 5)^4$. Find the Gradient and the Hessian and indicate what is special about the point $(1, 3, -5)$.

Solution

The gradient is defined as

$$\nabla f(x_1, x_2, x_3) = \begin{pmatrix} 4(x_1 - 1)^3 \\ 2(x_2 - 3) \\ 16(x_3 + 5)^3 \end{pmatrix}.$$

The Hessian is defined as

$$\nabla^2 f(x_1, x_2, x_3) = \begin{pmatrix} 12(x_1 - 1)^2 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 48(x_3 + 5)^2 \end{pmatrix}.$$

In the point $(1, 3, -5)$, the Gradient is $\nabla f(x_1, x_2, x_3) = (0, 0, 0)'$, and the Hessian equals

$$\nabla^2 f(x_1, x_2, x_3) = \begin{pmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{pmatrix}.$$

In the direction of x_1 and x_3 , the function surface is almost flat.

4 From non-iterative to iterative procedures

No exercises.

5 Least squares

TO DO.

6 Iteration-based Function Optimization

Chapter 2 on motivating problems is the first chapter that actually entails exercises.

6.1 Exercises (6.5 in the notes)

1. Suppose for every individual in a small pre-clinical study, it has been recorded how many epileptic seizures are observed (outcome y) and whether the individual is receiving a standard treatment (covariate $x = 0$) or experimental medication (covariate $x = 1$). The data are:

Subject i	Treatment x	# Seizures y
1	1	12
2	1	15
3	1	17
4	0	8
5	0	11
6	0	5

A Poisson regression model is put forward for these data, with linear predictor $\theta_i = \beta_0 + \beta_1 x_i$. Starting from $\beta^{(0)} = (0, 0)'$, do the following: Derive the likelihood equations. Can they be solved analytically in this case? Perform the first five steps of the Newton-Raphson algorithm to find the maximum of the likelihood. Put your results in a table with as columns: Iteration number, current point, and log-likelihood value. Do the same for Fisher-scoring.

Solution

The Poisson model yields

$$Y \sim \text{Poisson}(\lambda), \text{ with } f(y|\theta, \phi) = \frac{e^{-\lambda} \lambda^y}{y!},$$

and thus the likelihood L and log-likelihood ℓ are defined as

$$\begin{aligned}
L &= \prod_{i=1}^6 \frac{e^{-\lambda} \lambda^{y_i}}{y_i!} = \frac{e^{-e^{(\beta_0 + \beta_1 x_i)}} e^{(\beta_0 + \beta_1 x_i) y_i}}{y_i!} \\
\ell &= \sum_{i=1}^6 y_i \log \lambda - \lambda - \log(y_i!) \\
&= \sum_{i=1}^6 y_i (\beta_0 + \beta_1 x_i) - e^{(\beta_0 + \beta_1 x_i)} - \log(y_i!).
\end{aligned}$$

Accordingly, the first-order partial derivatives are defined as

$$\begin{aligned}
\frac{\partial \ell}{\partial \beta_0} &= \sum_{i=1}^6 y_i - e^{(\beta_0 + \beta_1 x_i)}, \\
\frac{\partial \ell}{\partial \beta_1} &= \sum_{i=1}^6 x_i y_i - x_i e^{(\beta_0 + \beta_1 x_i)},
\end{aligned}$$

and hence the Gradient (i.e., Score equation) can be written as

$$\nabla \ell(\beta_0, \beta_1) = S(\theta) = \begin{pmatrix} \sum_{i=1}^6 y_i - e^{(\beta_0 + \beta_1 x_i)}, \\ \sum_{i=1}^6 x_i (y_i - e^{(\beta_0 + \beta_1 x_i)}) \end{pmatrix}$$

Additionally, the second-order partial derivatives are defined as

$$\begin{aligned}
\frac{\partial^2 \ell}{\partial \beta_0^2} &= \sum_{i=1}^6 -e^{(\beta_0 + \beta_1 x_i)}, \\
\frac{\partial^2 \ell}{\partial \beta_1^2} &= \sum_{i=1}^6 -x_i^2 e^{(\beta_0 + \beta_1 x_i)}, \\
\frac{\partial^2 \ell}{\partial \beta_0 \partial \beta_1} &= \sum_{i=1}^6 -x_i e^{(\beta_0 + \beta_1 x_i)},
\end{aligned}$$

such that the Hessian $\nabla^2 \ell(\beta_0, \beta_1)$ can be written as

$$\nabla^2 \ell(\beta_0, \beta_1) = \begin{pmatrix} \sum_{i=1}^6 -e^{(\beta_0 + \beta_1 x_i)} & \sum_{i=1}^6 -x_i e^{(\beta_0 + \beta_1 x_i)} \\ \sum_{i=1}^6 -x_i e^{(\beta_0 + \beta_1 x_i)} & \sum_{i=1}^6 -x_i^2 e^{(\beta_0 + \beta_1 x_i)} \end{pmatrix}.$$

Setting the first-order partial derivatives to zero and filling in the data yields

$$S(\theta) = \begin{pmatrix} 68 - 3e^{(\beta_0 + \beta_1)} - 3e^{(\beta_0)} \\ 44 - 3e^{(\beta_0 + \beta_1)} \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \end{pmatrix}.$$

Hence, we have

$$\begin{aligned} 44 - 3e^{(\beta_0 + \beta_1)} &= 0 \\ 3e^{(\beta_0 + \beta_1)} &= 44, \end{aligned}$$

and thus

$$\begin{aligned} 68 - 3e^{(\beta_0)} &= 44 \\ 3e^{(\beta_0)} &= 24 \\ e^{(\beta_0)} &= 8 \\ \beta_0 &= \log 8 \approx 2.0794. \end{aligned}$$

Filling this into the previous equation yields

$$\begin{aligned} 3e^{(\log 8 + \beta_1)} &= 44 \\ \log 44 - \log 3 - \log 8 &= \beta_1 \approx 0.6061. \end{aligned}$$

Newton-Raphson method

```
NR <- function(formula, data = NULL, start, n.iter) {

  X <- model.matrix(formula, data)
  Y <- model.frame(formula, data)[,1]

  loglikelihood <- function(X, Y, beta) {
    constant <- sapply(Y, function(y) sum(log(1:y))) |> sum()
    sum(y - X %*% beta - exp(X %*% beta) - constant)
  }

  score <- function(X, Y, beta) {
    t(X) %*% (Y - exp(X %*% beta))
  }

  hess <- function(X, beta) {
    - t(X) %*% diag(c(exp(X %*% beta))) %*% X
  }
}
```



```

out <- matrix(0, n.iter+1, ncol(X))
out[1, ] <- b <- start

logL <- numeric(n.iter+1)
logL[1] <- loglikelihood(X, Y, b)

for (i in (1:n.iter)+1) {
  b <- b - solve(hess(X, b)) %*% score(X, Y, b)
  out[i, ] <- b
  logL[i] <- loglikelihood(X, Y, b)
}
data.frame(iter = 0:n.iter,
           out,
           logL = logL)
}

x <- c(1, 1, 1, 0, 0, 0)
y <- c(12, 15, 17, 8, 11, 5)

NR(y ~ x, start = c(0,0), n.iter = 20) |>
  knitr::kable() |>
  kableExtra::kable_styling(bootstrap_options = c("striped", "hover"))

```

Fisher scoring

Note that in this case, the expected Hessian equals

$$-x_i' \frac{\partial \mu_i}{\partial \theta_i} \nu_i^{-1} \frac{\partial \mu_i}{\partial \theta_i} x_i.$$

Given that

$$\frac{\partial \mu_i}{\partial \theta_i} = \frac{\partial \mu_i}{\partial \theta_i} \left(\exp\{\theta_i\} \right) = \exp \theta_i,$$

and

$$\nu_i^{-1} = \frac{1}{\exp\{\theta_i\}},$$

it follows that

iter	X1	X2	logL
0	0.000000	0.000000	-623.7158
1	7.000000	6.666667	-2589080.4867
2	6.007295	6.6593886	-952918.2151
3	5.026981	6.6397490	-351004.0668
4	4.079450	6.5874061	-129568.7184
5	3.214784	6.4524137	-48104.1538
6	2.536096	6.1320304	-18133.0024
7	2.169495	5.5011536	-7106.9141
8	2.083377	4.5941106	-3051.0568
9	2.079449	3.6165031	-1558.0226
10	2.079442	2.6657840	-1007.2910
11	2.079442	1.7932829	-803.7918
12	2.079442	1.0983733	-729.4703
13	2.079442	0.7096305	-705.1191
14	2.079442	0.6113113	-700.2547
15	2.079442	0.6061492	-700.0115
16	2.079442	0.6061358	-700.0108
17	2.079442	0.6061358	-700.0108
18	2.079442	0.6061358	-700.0108
19	2.079442	0.6061358	-700.0108
20	2.079442	0.6061358	-700.0108

$$\frac{\partial \mu_i}{\partial \theta_i} \nu_i^{-1} \frac{\partial \mu_i}{\partial \theta_i} = \exp\{\theta_i\} = \exp\{X\beta\}.$$

Hence, for the expected Hessian, we have

$$\mathcal{H} = E\left(\frac{\partial^2 \ell}{\partial \beta \partial \beta'}\right) = X' \text{diag}(\exp X\beta) X,$$

which is equal to the Hessian matrix $H(\beta)$, and thus Fisher scoring and Newton-Raphson are equivalent in this case.

2. Assume the function

$$f(x_1, x_2) = 8x_1 + 12x_2 + x_1^2 - 2x_2^2.$$


Sketch the contour lines of $f(x_1, x_2)$, and find the stationary point of $f(x_1, x_2)$. Does this point correspond to a minimum, a maximum, or something else?

Solution

```
library(ggplot2)
library(dplyr)
library(purrr)

fx1x2 <- function(x1, x2) 8*x1 + 12*x2 + x1^2 - 2*x2^2

expand.grid(x1 = -300:200/10,
            x2 = -200:300/10) |>
  mutate(f = map2_dbl(x1, x2, fx1x2)) |>
  ggplot(aes(x = x1, y = x2, z = f)) +
  stat_contour_filled(bins = 30, show.legend = FALSE) +
  theme_minimal() +
  labs(x = expression(italic(X[1])),
       y = expression(italic(X[2])),
       title = "Contour plot")
```



The first- and second-order partial derivatives of $f(x_1, x_2)$ are given by

$$f(x_1, x_2) = 8x_1 + 12x_2 + x_1^2 - 2x_2^2,$$

$$\frac{\partial f}{\partial x_1} = 8 + 2x_1,$$

$$\frac{\partial f}{\partial x_2} = 12 - 4x_2,$$

$$\frac{\partial^2 f}{\partial x_1^2} = 2,$$

$$\frac{\partial^2 f}{\partial x_2^2} = -4,$$

$$\frac{\partial^2 f}{\partial x_1 \partial x_2} = 0.$$

The stationary point of $f(x_1, x_2)$ is $f(-4, 3)$, which is a saddle point.

3. Consider the function


$$f(x_1, x_2) = 100(x_2 - x_1^2)^2 + (1 - x_1)^2.$$

Show that $(1, 1)'$ is a local minimizer of this function. Also, starting from the point $x^{(0)} = (0, 0)'$, perform the first five steps of the steepest descent and the Newton-Raphson algorithm to minimize the function. Put your results in a table with as columns: iteration number, current point, function value and gradient.

Solution

```
fx1x2 <- function(x1, x2) 100*(x2 - x1^2)^2 + (1 - x1)^2

expand.grid(x1 = -100:200/100,
            x2 = -100:200/100) |>
mutate(f = map2_dbl(x1, x2, fx1x2)) |>
ggplot(aes(x = x1, y = x2, z = f)) +
stat_contour_filled(bins = 50, show.legend = FALSE) +
theme_minimal() +
labs(x = expression(italic(X[1])),
     y = expression(italic(X[2])),
     title = "Contour plot")
```



Showing that the point $(1, 1)'$ is a local minimizer can be done by plugging the $(1, 1)'$ into the Gradient, and checking whether the Gradient equals zero,

$$\begin{aligned}
f(x_1, x_2) &= 100(x_2 - x_1^2)^2 + (1 - x_1)^2, \\
\nabla f(x_1, x_2) &= \begin{pmatrix} -400x_1(x_2 - x_1^2) - 2(1 - x_1) \\ 200(x_2 - x_1^2) \end{pmatrix}, \\
\nabla f(1, 1) &= \begin{pmatrix} -400(1 - 1) - 2(1 - 1) \\ 200(1 - 1) \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \end{pmatrix},
\end{aligned}$$

which shows that $(1, 1)'$ is a local minimizer. Moreover, the Hessian matrix is defined by

$$\nabla^2 f(x_1, x_2) = \begin{pmatrix} 1200x_1^2 - 400x_2 + 2 & -400x_1 \\ -400x_1 & 200 \end{pmatrix}$$

Steepest-Descent

```

f <- function(x1, x2) 100*(x2 - x1^2)^2 + (1 - x1)^2

score <- function(x1, x2) {
  c(400*x1^3 - 400*x1*x2 + 2*x1 - 2,
    200*x2 - 200*x1^2)
}

hess <- function(x1, x2) {
  matrix(c(1200*x1^2 - 400*x2 + 2, -400*x1, -400*x1, 200),
    nrow = 2, ncol = 2)
}

SD <- function(start, n.iter, alpha, rho, tol = 1e-16) {
  b <- start
  grad <- matrix(0, n.iter + 1, 2)
  grad[1, ] <- score(b[1], b[2])

  out <- matrix(0, n.iter + 1, 2)
  out[1, ] <- b

  i <- 1; conv <- FALSE

  while (!conv) {
    i <- i+1
    fold <- f(out[i-1, 1], out[i-1, 2])
    gradvec <- score(out[i-1, 1], out[i-1, 2])
    out[i, ] <- out[i-1, ] - alpha * gradvec / sum(gradvec^2)
    grad[i, ] <- gradvec
  }
}

```

```

fnew <- f(out[i,1], out[i,2])
a <- alpha
while(fnew > fold) {
  a <- a*rho
  out[i, ] <- out[i-1, ] - a * gradvec / sum(gradvec^2)
  grad[i, ] <- c(score(out[i,1], out[i,2]))
  fnew <- f(out[i,1], out[i,2])
}
if (
  i - 1 == n.iter |
  abs(fnew - fold) < tol
) {
  conv <- TRUE
}

}
data.frame(iter = 0:(nrow(out)-1),
           out = out,
           grad = grad,
           fval = f(out[,1], out[,2])) |>
  subset(iter < i)
}

SDout <- SD(c(0,0), 20000, 1, 0.8, 1e-10)

SDout |>
  head(15) |>
  knitr::kable() |>
  kableExtra::kable_styling(bootstrap_options = c("striped", "hover"))

```

Newton-Raphson

```

NR <- function(start, n.iter, alpha, rho) {

  b <- start
  grad <- matrix(0, n.iter + 1, 2)
  grad[1, ] <- score(b[1], b[2])

  out <- matrix(0, n.iter + 1, 2)
  out[1, ] <- b

```

iter	out.1	out.2	grad.1	grad.2	fval
0	0.0000000	0.0000000	-2.0000000	0.0000000	1.0000000
1	0.2560000	0.0000000	5.2228864	-13.107200	0.9830327
2	0.2297645	0.0658398	5.2228864	-13.107200	0.6102878
3	0.2503131	0.0462667	0.1416746	-3.277992	0.5888936
4	0.2491826	0.0724227	-2.5313281	2.066142	0.5743991
5	0.2695487	0.0557993	0.3566254	-3.371428	0.5619755
6	0.2668834	0.0809960	-2.5091373	1.953857	0.5470039
7	0.2881952	0.0644006	0.7270034	-3.731174	0.5414702
8	0.2827931	0.0921256	-2.8092036	2.430734	0.5291569
9	0.3046506	0.0732128	0.9976613	-3.919835	0.5219235
10	0.2981029	0.0989389	-2.6049709	2.014701	0.5028071
11	0.3187362	0.0829810	1.0103687	-3.722352	0.4987602
12	0.3114437	0.1098474	-2.9779538	2.570033	0.4906224
13	0.3321087	0.0920131	1.0930160	-3.656631	0.4795061
14	0.3240513	0.1189688	-3.1613443	2.791914	0.4763936

```

for (i in 1:n.iter + 1) {

  fold <- f(out[i-1,1], out[i-1,2])
  b <- out[i - 1, ]

  out[i, ] <- b - solve(hess(b[1], b[2])) %*% score(b[1], b[2])
  grad[i, ] <- c(score(b[1], b[2]))

  fnew <- f(out[i,1], out[i,2])
  a <- alpha
  while (fnew>fold){
    a <- a*rho
    out[i,] <- out[i-1,] - a * solve(hess(b[1], b[2])) %*% score(b[1], b[2])
    grad[i, ] <- c(score(out[i,1], out[i,2]))
    fnew <- f(out[i,1], out[i,2])
  }
}

data.frame(iter = 0:(nrow(out)-1),
           out = out,
           grad = grad,
           fval = f(out[,1], out[,2]))
}

NRout <- NR(c(0,0), 15, 1, 0.8)

```

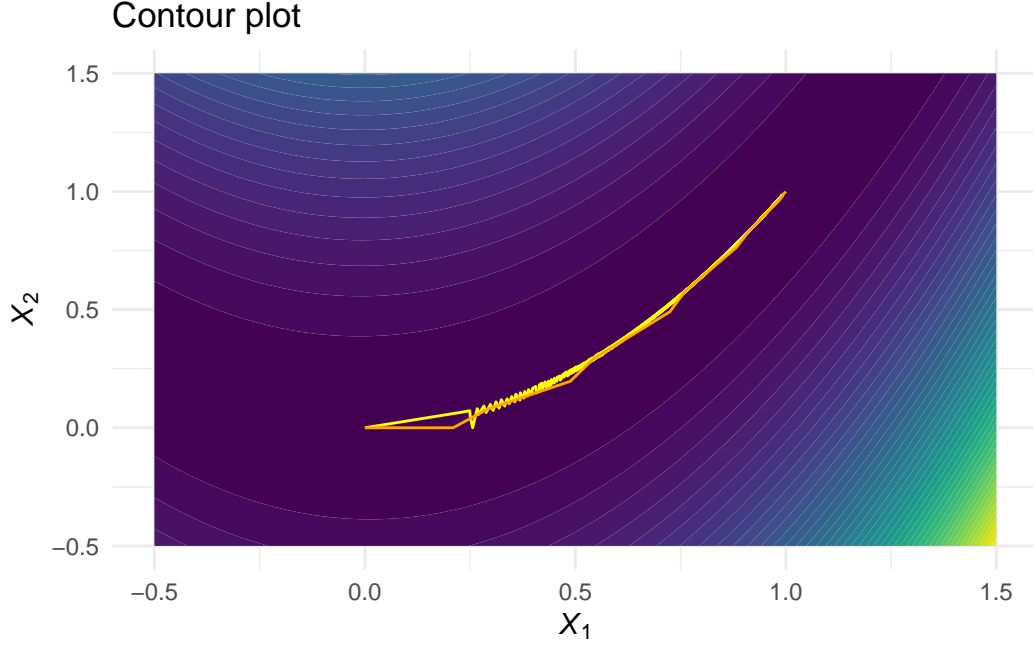

iter	out.1	out.2	grad.1	grad.2	fval
0	0.0000000	0.0000000	-2.0000000	0.0000000	1.0000000
1	0.2097152	0.0000000	2.1087792	-8.7960930	0.8179782
2	0.2903887	0.0778174	2.1087792	-8.7960930	0.5077839
3	0.4877049	0.1965794	7.0277391	-8.2553296	0.4328224
4	0.5430563	0.2918463	7.0277391	-8.2553296	0.2097363
5	0.7243882	0.4907541	9.2958913	-6.7968490	0.1914547
6	0.7597374	0.5759513	9.2958913	-6.7968490	0.0578823
7	0.8827605	0.7636815	5.2684843	-3.1169062	0.0380329
8	0.9112381	0.8295438	5.2684843	-3.1169062	0.0079444
9	0.9876125	0.9695454	0.1180717	-0.1621945	0.0035559
10	0.9933299	0.9866717	2.2795435	-1.1666107	0.0000446
11	0.9999567	0.9998694	-0.0003516	-0.0065379	0.0000002
12	0.9999996	0.9999992	0.0174780	-0.0087827	0.0000000
13	1.0000000	1.0000000	0.0000000	-0.0000004	0.0000000
14	1.0000000	1.0000000	0.0000000	0.0000000	0.0000000
15	1.0000000	1.0000000	0.0000000	0.0000000	0.0000000

```

NRout |>
  knitr::kable() |>
  kableExtra::kable_styling(bootstrap_options = c("striped", "hover"))

expand.grid(x1 = -50:150/100,
            x2 = -50:150/100) |>
  mutate(f = map2_dbl(x1, x2, fx1x2)) |>
  ggplot(aes(x = x1, y = x2, z = f)) +
  stat_contour_filled(bins = 50, show.legend = FALSE) +
  geom_line(data = SDout,
            mapping = aes(x = out.1, y = out.2, z = NULL),
            col = "yellow") +
  geom_line(data = NRout,
            mapping = aes(x = out.1, y = out.2, z = NULL),
            col = "orange") +
  theme_minimal() +
  labs(x = expression(italic(X[1])),
       y = expression(italic(X[2])),
       title = "Contour plot")

```



4. Suppose for an individual during consecutive nights, it is recorded how loudly he snores (covariate x) and whether he wakes up or not (the outcome Y). Consider the following hypothetical data are collected: $x = (0, 1, 2, 3, 4, 5)'$ and $y = (0, 1, 0, 1, 1, 1)'$. A logistic regression model is put forward for these data such that $\text{logit}(\Pr(y_i = 1|x_i)) = \text{logit}(\pi(x_i)) = \beta_0 + \beta_1 x_i$. Starting from $\beta^{(0)} = (0, 0)'$, perform the first five steps of the Newton-Raphson algorithm to find the maximum of the likelihood. Put your results in a table with as columns: iteration number, current point and loglikelihood value. Do the same for iterative reweighted least squares.

Solution

For logistic regression, the likelihood is defined as

$$L = \prod_{i=1}^N \frac{e^{y_i(\beta_0 + \beta_1 x_i)}}{1 + e^{(\beta_0 + \beta_1 x_i)}},$$

$$\ell = \sum_{i=1}^N y_i(\beta_0 + \beta_1 x_i) - \log(1 + e^{(\beta_0 + \beta_1 x_i)}).$$

Additionally, the Gradient is defined by

$$\nabla \ell(\beta_0, \beta_1) = \begin{pmatrix} \sum_{i=1}^N y_i - \frac{e^{(\beta_0 + \beta_1 x_i)}}{1 + e^{(\beta_0 + \beta_1 x_i)}} \\ \sum_{i=1}^N x_i (y_i - \frac{e^{(\beta_0 + \beta_1 x_i)}}{1 + e^{(\beta_0 + \beta_1 x_i)}}) \end{pmatrix},$$

while the Hessian is defined as

$$\nabla^2 \ell(\beta_0, \beta_1) = \begin{pmatrix} -\sum_{i=1}^N \frac{e^{(\beta_0 + \beta_1 x_i)}}{(1 + e^{(\beta_0 + \beta_1 x_i)})^2} & -\sum_{i=1}^N x_i \frac{e^{(\beta_0 + \beta_1 x_i)}}{(1 + e^{(\beta_0 + \beta_1 x_i)})^2} \\ -\sum_{i=1}^N x_i \frac{e^{(\beta_0 + \beta_1 x_i)}}{(1 + e^{(\beta_0 + \beta_1 x_i)})^2} & -\sum_{i=1}^N x_i^2 \frac{e^{(\beta_0 + \beta_1 x_i)}}{(1 + e^{(\beta_0 + \beta_1 x_i)})^2} \end{pmatrix}.$$

```
loglikelihood <- function(X, Y, beta) {
  sum(Y * (X%*%beta) - log(1 + exp(X %*% beta)))
}
score <- function(X, Y, beta) {
  t(X) %*% (Y - 1/(1 + exp(-X%*%beta)))
}
hess <- function(X, Y, beta) {
  - t(X) %*% diag(c(exp(-(X%*%beta))/(1 + exp(-X%*%beta))^2)) %*% X
}
```

Newton-Raphson implementation

```
NRlogistic <- function(formula, data = NULL, start, n.iter) {
  X <- model.matrix(formula, data)
  Y <- model.frame(formula, data)[, 1]

  out <- matrix(0, n.iter+1, ncol(X))
  out[1, ] <- b <- start

  logL <- numeric(n.iter+1)
  logL <- loglikelihood(X, Y, b)

  for (i in 1:n.iter + 1) {
    b <- b - solve(hess(X, Y, b)) %*% score(X, Y, b)
    out[i, ] <- b
    logL[i] <- loglikelihood(X, Y, b)
  }

  data.frame(iter = 0:n.iter,
             b0 = out[,1],
             b1 = out[,2],
             logL = logL)
}

x <- c(0,1,2,3,4,5)
y <- c(0,1,0,1,1,1)
```

iter	b0	b1	logL
0	0.000000	0.000000	-4.158883
1	-1.047619	0.6857143	-2.626827
2	-1.444172	0.9933894	-2.457094
3	-1.602433	1.1249532	-2.440395
4	-1.624928	1.1443026	-2.440125
5	-1.625338	1.1446616	-2.440125

```
NRlogistic(y ~ x, start = c(0,0), n.iter = 5) |>
  knitr::kable() |>
  kableExtra::kable_styling(bootstrap_options = c("striped", "hover"))
```

```
glm(y ~ x, family = binomial) |> summary()
```

Call:

```
glm(formula = y ~ x, family = binomial)
```

Deviance Residuals:

1	2	3	4	5	6
-0.5995	1.3872	-1.4692	0.5509	0.3189	0.1815

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.6253	1.9284	-0.843	0.399
x	1.1447	0.9278	1.234	0.217

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 7.6382 on 5 degrees of freedom
 Residual deviance: 4.8802 on 4 degrees of freedom
 AIC: 8.8802

Number of Fisher Scoring iterations: 5

Convergence is reached after five iterations!

Iterative re-weighted least squares implementation

iter	b0	b1	logL
0	0.000000	0.000000	-4.158883
1	-1.047619	0.6857143	-2.626827
2	-1.444172	0.9933894	-2.457094
3	-1.602433	1.1249532	-2.440395
4	-1.624928	1.1443026	-2.440125
5	-1.625338	1.1446616	-2.440125

```
IRLS <- function(formula, data = NULL, start, n.iter) {

  X <- model.matrix(formula, data)
  Y <- model.frame(formula, data)[,1]
  out <- matrix(0, n.iter+1, ncol(X))
  out[1, ] <- b <- start

  logL <- numeric(n.iter+1)
  logL[1] <- loglikelihood(X, Y, b)

  for (i in 1:n.iter+1) {
    e <- exp(X%*%b) / (1 + exp(X%*%b))
    W <- diag(c(e / (1+exp(X %*% b))))
    Z <- X %*% b + (y - e) * (1 / (e*(1-e)))
    b <- solve(t(X) %*% W %*% X) %*% t(X) %*% W %*% Z
    out[i, ] <- b
    logL[i] <- loglikelihood(X, Y, b)
  }
  data.frame(iter = 0:n.iter,
             b0 = out[,1],
             b1 = out[,2],
             logL = logL)
}

IRLS(y ~ x, start = c(0,0), n.iter = 5) |>
knitr::kable() |>
kableExtra::kable_styling(bootstrap_options = c("striped", "hover"))
```

And again, convergence is reached after five iterations!

5. Consider the function

$$f(x) = \frac{e^x}{(1 + e^x)^2}.$$

Using an iterative procedure of your liking, find the optimum of the function, and check whether it is a minimum or a maximum.

Solution

First, calculate we calculate the derivatives.

$$\begin{aligned}f(x) &= \frac{e^x}{(1 + e^x)^2}, \\f'(x) &= \frac{e^x - e^{2x}}{(1 + e^x)^3}, \\f''(x) &= \frac{e^x - 4e^{2x} + e^{3x}}{(1 + e^x)^4}.\end{aligned}$$

We can first find the optimum analytically. Let's first take the log of the function, which makes it easier to work with:

$$\log f(x) = x - 2\log(1 + e^x).$$

Subsequently, we take the derivative of $\log f(x)$ and set it equal to zero to find the optimum.

$$\begin{aligned}\frac{\partial f}{\partial x} &= 1 - \frac{2e^x}{1 + e^x} = 0, \\ \implies \frac{2e^x}{1 + e^x} &= 1, \\ 1 + e^x &= 2e^x, \\ e^x &= 1, \\ x &= 0.\end{aligned}$$

So we know the solution must be $x = 0$. Doing the same steps using the Newton-Raphson algorithm yields

```
fx <- function(x) -exp(x) / (1 + exp(x))^2
f1x <- function(x) -(exp(x) - exp(2*x)) / (1 + exp(x))^3
f2x <- function(x) (exp(x) - 4*exp(2*x) + exp(3*x)) / (1 + exp(x))^4

NR <- function(start = 0.5, n.iter = 20, alpha = 1, rho = 0.8) {
  out <- matrix(0, n.iter+1, 4)
  out[1, ] <- c(start, fx(start), f1x(start), f2x(start))

  colnames(out) <- c("x", "fx", "f1x", "f2x")
}
```

x	fx	f1x	f2x
0.5	-0.2350037	0.0575568	-0.0963568
0.5	-0.2350037	0.0575568	-0.0963568
0.5	-0.2350037	0.0575568	-0.0963568
0.5	-0.2350037	0.0575568	-0.0963568
0.5	-0.2350037	0.0575568	-0.0963568
0.5	-0.2350037	0.0575568	-0.0963568

```

for (i in 1:n.iter + 1) {
  a <- alpha
  new <- out[i-1, 1] - a * out[i-1, 3] / out[i-1, 4]
  out[i, ] <- c(new, fx(new), f1x(new), f2x(new))
  while(out[i, 2] > out[i-1, 2]) {
    a <- a*rho
    new <- out[i-1, 1] - a * out[i-1, 3] / out[i-1, 4]
    out[i, ] <- c(new, fx(new), f1x(new), f2x(new))
  }
}
out
}

NR(0.5, 5) |>
knitr::kable() |>
kableExtra::kable_styling(bootstrap_options = c("striped", "hover"))

```

6. Continuation of exercise 6 from chapter 3: implement maximum likelihood estimation for this logistic regression.

Solution

```

x <- c(0.5, 1, 1.5, 2, 2.5)
y <- c(0,0,1,0,1)

NRlogistic(y ~ x, start = c(0,0), n.iter = 5) |>
knitr::kable() |>
kableExtra::kable_styling(bootstrap_options = c("striped", "hover"))

IRLS(y ~ x, start = c(0,0), n.iter = 5) |>
knitr::kable() |>
kableExtra::kable_styling(bootstrap_options = c("striped", "hover"))

```

iter	b0	b1	logL
0	0.000000	0.000000	-3.465736
1	-2.800000	1.600000	-2.479523
2	-3.698907	2.079500	-2.423599
3	-3.886773	2.177155	-2.421969
4	-3.893957	2.180846	-2.421967
5	-3.893967	2.180851	-2.421967

iter	b0	b1	logL
0	0.000000	0.000000	-3.465736
1	-2.800000	1.600000	-2.479523
2	-3.698907	2.079500	-2.423599
3	-3.886773	2.177155	-2.421969
4	-3.893957	2.180846	-2.421967
5	-3.893967	2.180851	-2.421967

```
glm(y ~ x, family = binomial) |> summary()
```

Call:

```
glm(formula = y ~ x, family = binomial)
```

Deviance Residuals:

1	2	3	4	5
-0.3430	-0.5758	1.4506	-1.3814	0.6181

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-3.894	3.465	-1.124	0.261
x	2.181	1.950	1.119	0.263

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 6.7301 on 4 degrees of freedom
Residual deviance: 4.8439 on 3 degrees of freedom
AIC: 8.8439

Number of Fisher Scoring iterations: 4

References